

# Knowledge Tracing Machines: Factorization Machines for Knowledge Tracing

Jill-Jênn Vie    Hisashi Kashima



京都大学  
KYOTO UNIVERSITY

The University of Hawai'i at Mānoa, January 24, 2019

<https://arxiv.org/abs/1811.03388>

# Who is this person?



... Actually, she does not exist



2014



2015



2016



2017



2018

5 years of GAN progress on face generation

A Style-Based Generator Architecture for Generative Adversarial Networks (Karras, Laine and Aila, 2018)

[arxiv.org/abs/1812.04948](https://arxiv.org/abs/1812.04948)

# AI for Social Good

AI can:

- generate fakes
- recognize images/speech
- predict the next word
- play go (make decisions)

as long as you have enough data.

Can it also:

- generate exercises
- improve education
- predict student performance
- optimize human learning

as long as you have enough data?

My research

Use machine learning techniques on data that comes from humans

# Outline

- ① Introduction to machine learning
- ② Knowledge Tracing and existing models
- ③ Knowledge Tracing Machines
  - Encoding data into sparse features
  - Running logistic regression or factorization machines
- ④ Experiments and results

# What is machine learning?

Train a model on some data (e.g. pictures of cats)

Test it on new data

Tweak the parameters to optimize a value (e.g., maximize accuracy)

## How to choose the best model?

Train model on some of your data, keep the remaining data hidden

Try to predict the remaining data

## What is a good machine learning model?

Can reconstruct training data at 90%



predicts new points correctly 80% of the time

## What is a bad machine learning model?

Can reconstruct training data at 100%

predicts new points correctly 20% of the time


# It is called overfitting



Underfitting



Desired



Overfitting

# Formally

We have existing data  $(\mathbf{x}, y)$   
and we want to learn a function  $y = f(\mathbf{x})$ .

## Examples

**classification**  $\mathbf{x}$  can be an image and  $y \in \{man, woman\}$

**regression**  $\mathbf{x}$  can be the temperature over the last 7 days  
and  $y$  is the temperature of the next day



# Two typical examples


For regression: **linear regression**

We have existing data  $(\mathbf{x}, y) \in \mathbb{R}^d \times \mathbb{R}$  and we want to learn a function  $y = \langle \mathbf{w}, \mathbf{x} \rangle + b = \sum_i w_i x_i + b$ .

Example if  $d = 1$ : we want to learn  $w x + b = y$ , an affine function!

For classification: **logistic regression**

We have existing data  $(\mathbf{x}, y) \in \mathbb{R}^d \times \{0, 1\}$  and we want to learn a function  $y = \sigma(\langle \mathbf{w}, \mathbf{x} \rangle + b)$ .



# Is accuracy a good metric for classification?

If 99% of people are in good health, and you are making a machine learning algorithm that **always** predicts a person is healthy...  
It will have accuracy 99%.

Is this a good machine?

Better metrics: AUC, F1 score, etc.

# Students try exercises

## Math Learning

Items	$5 - 5 = ?$	$17 - 3 = ?$	$13 - 7 = ?$
New student	○	○	×

## Language Learning

	PRON	VERB	PRON	NOUN	CONJ	PRON	VERB	PRON	NOUN
correct:	She	is	my	mother	and	he	is	my	father
student:	she	is		mader	and	he	is		fhader
label:	○	○	×	×	○	○	○	×	×

## Challenges

- Users can attempt a same item multiple times
- Users learn over time
- People can make mistakes that do not reflect their knowledge

# Predicting student performance: knowledge tracing

## Data

A population of users answering items

- Events: “User  $i$  answered item  $j$  correctly/incorrectly”

Side information

- We know the skills required to solve each item e.g., +, ×
- Class ID, school ID, etc.

## Goal: classification problem

Predict the performance of new users on existing items

## Method

Learn parameters of questions from historical data e.g., *difficulty*

Measure parameters of new students e.g., *expertise*

# Existing work

Model	Basically	Original AUC	Fixed AUC
Bayesian Knowledge Tracing (Corbett and Anderson 1994)	Hidden Markov Model	0.67	0.63
Deep Knowledge Tracing (Piech et al. 2015)	Recurrent Neural Network	0.86	0.75
Item Response Theory (Rasch 1960) (Wilson et al., 2016)	Online Logistic Regression		0.76

$$\underbrace{\text{PFA}}_{\text{LogReg}} \leq \underbrace{\text{DKT}}_{\text{LSTM}} \leq \underbrace{\text{IRT}}_{\text{LogReg}} \leq \underbrace{\text{KTM}}_{\text{FM}}$$

# Limitations and contributions

- Several models for knowledge tracing were developed independently
- In our paper, we prove that our approach is more generic

## In this paper

- Knowledge Tracing Machines unify most existing models
  - Encoding student data to sparse features
  - Then running logistic regression or factorization machines
- Better models found
  - It is better to estimate a bias per item, not only per skill
  - Side information improves performance more than higher dim.

# Our small dataset

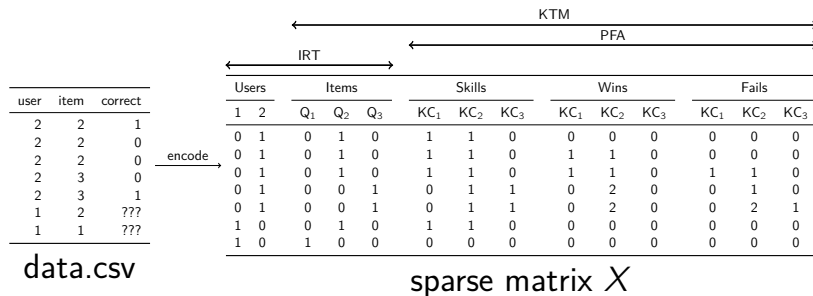
- User 1 answered Item 1 correct
- User 1 answered Item 2 incorrect
- User 2 answered Item 1 incorrect
- User 2 answered Item 1 correct
- User 2 answered Item 2 ???

user	item	correct
1	1	1
1	2	0
2	1	0
2	1	1
2	2	???

`dummy.csv`

# Our approach

- Encode data to sparse features



- Run logistic regression or factorization machines



# Model 1: Item Response Theory

Learn abilities  $\theta_i$  for each user  $i$

Learn easiness  $e_j$  for each item  $j$  such that:

$$Pr(\text{User } i \text{ Item } j \text{ OK}) = \sigma(\theta_i + e_j) \quad \sigma : x \mapsto 1/(1 + \exp(-x))$$

$$\text{logit } Pr(\text{User } i \text{ Item } j \text{ OK}) = \theta_i + e_j$$


Really popular model, used for the PISA assessment

## Logistic regression

Learn  $\mathbf{w}$  such that  $\text{logit } Pr(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle + b$

# Graphically: IRT as logistic regression

Encoding “User  $i$  answered Item  $j$ ” with **sparse features**:



$$\langle \mathbf{w}, \mathbf{x} \rangle = \theta_i + e_j = \text{logit } Pr(\text{User } i \text{ Item } j \text{ OK})$$

# Encoding into sparse features

Users			Items		
$U_0$	$U_1$	$U_2$	$I_0$	$I_1$	$I_2$
0	1	0	0	1	0
0	1	0	0	0	1
0	0	1	0	1	0
0	0	1	0	1	0
0	0	1	0	0	1

Then logistic regression can be run on the sparse features.

# Oh, there's a problem

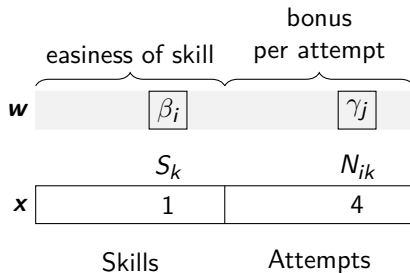
	Users			Items			$y_{\text{pred}}$	$y$
	$U_0$	$U_1$	$U_2$	$I_0$	$I_1$	$I_2$		
User 1 Item 1 OK	0	1	0	0	1	0	0.575135	1
User 1 Item 2 NOK	0	1	0	0	0	1	0.395036	0
User 2 Item 1 <b>NOK</b>	0	0	1	0	1	0	<b>0.545417</b>	<b>0</b>
User 2 Item 1 <b>OK</b>	0	0	1	0	1	0	<b>0.545417</b>	<b>1</b>
User 2 Item 2 NOK	0	0	1	0	0	1	0.366595	0

We predict the same thing when there are several attempts.

# Count number of attempts: AFM

Keep a counter of attempts at skill level:

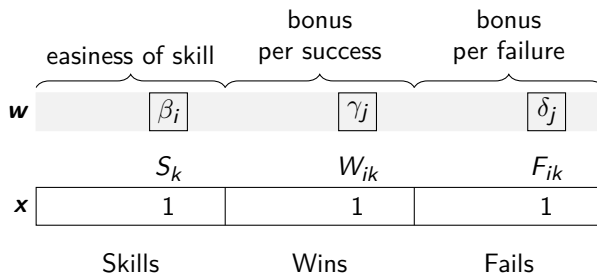
user	item	skill	correct	attempts (for the same skill)
1	1	1	1	0
1	2	2	0	0
2	1	1	0	0
2	1	1	1	1
2	2	2	0	0



# Count successes and failures: PFA

Count separately successes  $W_{ik}$  and fails  $F_{ik}$  of student  $i$  over skill  $k$ .

user	item	skill	correct	wins	fails
1	1	1	1	0	0
1	2	2	0	0	0
2	1	1	0	0	0
2	1	1	1	0	1
2	2	2	0	0	0



## Model 2: Performance Factor Analysis

$W_{ik}$ : how many successes of user  $i$  over skill  $k$  ( $F_{ik}$ : #failures)

Learn  $\beta_k$ ,  $\gamma_k$ ,  $\delta_k$  for each skill  $k$  such that:

$$\text{logit } Pr(\text{User } i \text{ Item } j \text{ OK}) = \sum_{\text{Skill } k \text{ of Item } j} \beta_k + W_{ik}\gamma_k + F_{ik}\delta_k$$

Skills			Wins			Fails		
$S_0$	$S_1$	$S_2$	$S_0$	$S_1$	$S_2$	$S_0$	$S_1$	$S_2$
0	1	0	0	0	0	0	0	0
0	0	1	0	0	0	0	0	0
0	1	0	0	0	0	0	0	0
0	1	0	0	0	0	0	1	0
0	0	1	0	0	0	0	0	0

[illegible]



# Test on a large dataset: Assistments 2009

346860 attempts of 4217 students over 26688 items on 123 skills.

model	dim	AUC	improvement
PFA: skills, wins, fails	0	0.685	+0.07
AFM: skills, attempts	0	0.616	

## Model 3: a new model (but still logistic regression)

model	dim	AUC	improvement
KTM: items, skills, wins, fails	0	0.746	+0.06
IRT: users, items	0	0.691	
PFA: skills, wins, fails	0	0.685	+0.07
AFM: skills, attempts	0	0.616	

# Here comes a new challenger

How to model **pairwise interactions** with **side information**?

## Logistic Regression

Learn a 1-dim **bias** for each feature (each user, item, etc.)

## Factorization Machines

Learn a 1-dim **bias** and a  $k$ -dim **embedding** for each feature

# How to model pairwise interactions with side information?

If you know user  $i$  attempted item  $j$  on **mobile** (not desktop)

How to model it?

$y$ : score of event “user  $i$  solves correctly item  $j$ ”

IRT

$$y = \theta_i + e_j$$


Multidimensional IRT (similar to collaborative filtering)

$$y = \theta_i + e_j + \langle \mathbf{v}_{\text{user } i}, \mathbf{v}_{\text{item } j} \rangle$$


With side information

$$y = \theta_i + e_j + w_{\text{mobile}} + \langle \mathbf{v}_{\text{user } i}, \mathbf{v}_{\text{item } j} \rangle + \langle \mathbf{v}_{\text{user } i}, \mathbf{v}_{\text{mobile}} \rangle + \langle \mathbf{v}_{\text{item } j}, \mathbf{v}_{\text{mobile}} \rangle$$

# Graphically: logistic regression




# Graphically: factorization machines



# Formally: factorization machines

Each **user**, **item**, **skill**  $k$  is modeled by bias  $w_k$  and embedding  $v_k$ .



$$\text{logit } p(\mathbf{x}) = \mu + \underbrace{\sum_{k=1}^N w_k x_k}_{\text{logistic regression}} + \underbrace{\sum_{1 \leq k < l \leq N} x_k x_l \langle \mathbf{v}_k, \mathbf{v}_l \rangle}_{\text{pairwise relationships}}$$

Steffen Rendle (2012). “Factorization Machines with libFM”. In: *ACM Transactions on Intelligent Systems and Technology (TIST)* 3.3, 57:1–57:22. DOI: 10.1145/2168752.2168771

# Training using MCMC

Priors:  $w_k \sim \mathcal{N}(\mu_0, 1/\lambda_0)$     $\mathbf{v}_k \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Lambda}^{-1})$

Hyperpriors:  $\mu_0, \dots, \mu_n \sim \mathcal{N}(0, 1), \lambda_0, \dots, \lambda_n \sim \Gamma(1, 1) = U(0, 1)$

---

**Algorithm 1** MCMC implementation of FMs

---

**for** each iteration **do**

    Sample hyperp.  $(\lambda_i, \mu_i)_i$  from posterior using Gibbs sampling

    Sample weights  $\mathbf{w}$

    Sample vectors  $\mathbf{V}$

    Sample predictions  $\mathbf{y}$

**end for**

---

Implementation in C++ (libFM) with Python wrapper (pyWFM).

Steffen Rendle (2012). “Factorization Machines with libFM”. In: *ACM Transactions on Intelligent Systems and Technology (TIST)* 3.3, 57:1–57:22. DOI: 10.1145/2168752.2168771




# Datasets


Name	Users	Items	Skills	Skills/i	Entries	Sparsity	Attempts/u
fraction	536	20	8	2.800	10720	0.000	1.000
timss	757	23	13	1.652	17411	0.000	1.000
ecpe	2922	28	3	1.321	81816	0.000	1.000
assistments	4217	26688	123	0.796	346860	0.997	1.014
berkeley	1730	234	29	1.000	562201	0.269	1.901
castor	58939	17	2	1.471	1001963	0.000	1.000

# Accuracy results on the Assistments dataset


Effect of data



Effect of dimension




# AUC results on the Assistments dataset



model	dim	AUC	improvement
KTM: items, skills, wins, fails, extra	5	0.819	
KTM: items, skills, wins, fails, extra	0	0.815	+0.05
KTM: items, skills, wins, fails	10	0.767	
KTM: items, skills, wins, fails	0	0.759	+0.02
DKT (Wilson et al., 2016)	100	0.743	+0.05
IRT: users, items	0	0.691	
PFA: skills, wins, fails	0	0.685	+0.07
AFM: skills, attempts	0	0.616	

# Bonus: interpreting the learned embeddings



# Take home message

**Knowledge tracing machines** unify many existing EDM models

- Side information improves performance more than higher  $d$
- We can visualize learning (and provide feedback to learners)

It has better results than **deep neural networks**

- It is important to find simple models that fit data
- Instead of monster models that overfit data

# Do you have any questions?

Read our article:

Knowledge Tracing Machines

<https://arxiv.org/abs/1811.03388>

Try the code:

<https://github.com/jilljenn/ktm>

Feel free to chat (or tell me about your favorite manga or anime):

[vie@jill-jenn.net](mailto:vie@jill-jenn.net)



Corbett, Albert T and John R Anderson (1994). “Knowledge tracing: Modeling the acquisition of procedural knowledge”. In: *User modeling and user-adapted interaction* 4.4, pp. 253–278.



Piech, Chris, Jonathan Bassen, Jonathan Huang, Surya Ganguli, Mehran Sahami, Leonidas J Guibas, and Jascha Sohl-Dickstein (2015). “Deep knowledge tracing”. In: *Advances in Neural Information Processing Systems (NIPS)*, pp. 505–513.



Rasch, Georg (1960). “Studies in mathematical psychology: I. Probabilistic models for some intelligence and attainment tests.”. In:



Rendle, Steffen (2012). “Factorization Machines with libFM”. In: *ACM Transactions on Intelligent Systems and Technology (TIST)* 3.3, 57:1–57:22. DOI: 10.1145/2168752.2168771.